

# Heart Disease Risk Assessment - Technical Architecture

## System Overview

This document outlines the technical architecture of a production-grade machine learning system for cardiovascular risk assessment. The system employs modern software engineering practices with microservices architecture, containerization, and automated deployment.

## Architecture Patterns

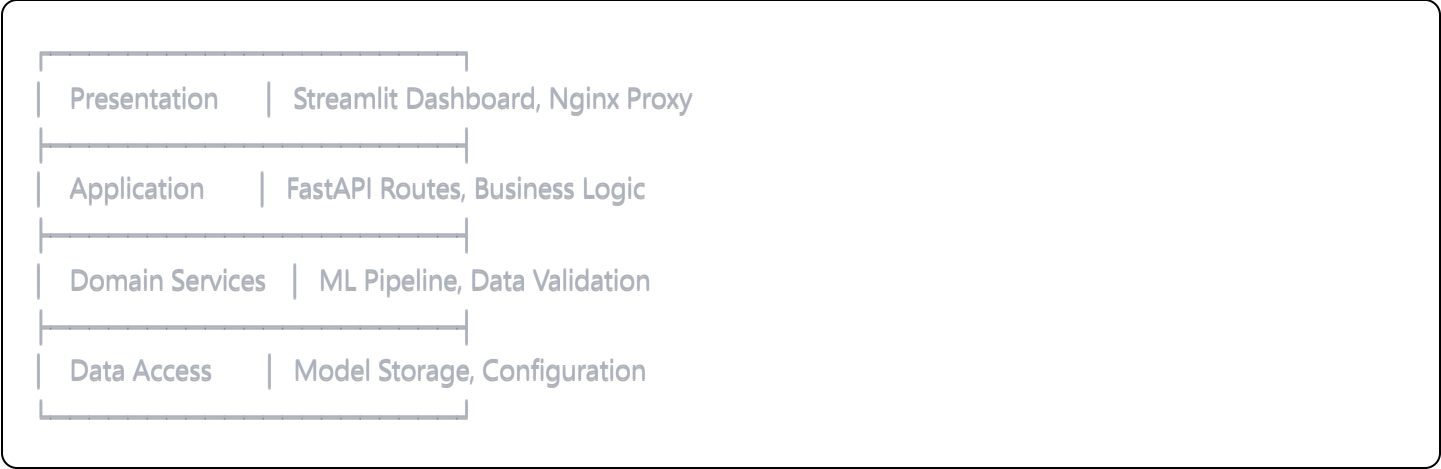
### Microservices Architecture

- **API Service:** FastAPI backend for ML predictions and data processing
- **Dashboard Service:** Streamlit frontend for user interaction
- **Proxy Service:** Nginx reverse proxy for load balancing and SSL termination
- **Service Communication:** RESTful APIs with JSON data exchange

### Event-Driven Design

- Asynchronous processing for batch operations
- Real-time health monitoring and alerting
- Scalable architecture supporting concurrent users

### Layered Architecture



## Technology Stack

### Machine Learning Pipeline

### Framework and Libraries

- **Core ML Framework:** scikit-learn 1.3+ for model training and inference
- **Ensemble Methods:** Random Forest, XGBoost, Logistic Regression
- **Feature Engineering:** Custom medical domain transformations
- **Model Interpretation:** SHAP 0.42+ for explainable AI
- **Data Processing:** pandas 2.0+, NumPy 1.24+ for data manipulation

## Model Architecture

Input Features (13) → Feature Engineering (20) → Ensemble Models → Risk Assessment

├ Random Forest  
├ XGBoost  
└ Logistic Regression

## Performance Specifications

- **Training Dataset:** UCI Heart Disease (303 patients)
- **Cross-Validation:** Stratified 5-fold with 81.39% average accuracy
- **Test Performance:** 86.89% accuracy, 95.35% AUC-ROC
- **Feature Engineering:** 20 features from 13 base clinical parameters

## Backend Services

### API Framework

- **Framework:** FastAPI 0.101+ for high-performance async REST services
- **Data Validation:** Pydantic 2.1+ for request/response schemas
- **Authentication:** JWT token framework (ready for implementation)
- **Documentation:** Automatic OpenAPI/Swagger generation
- **ASGI Server:** Uvicorn for production deployment

## Service Architecture

mermaid

graph TD

A[Client Request] --> B[Nginx Reverse Proxy]

B --> C[FastAPI Application]

C --> D[Request Validation]

D --> E[Business Logic Layer]

E --> F[ML Prediction Service]

F --> G[Data Preprocessing]

G --> H[Model Inference]

H --> I[Result Interpretation]

I --> J[Response Formatting]

J --> C

## Prediction Service Components

- **Data Validator:** Medical parameter range checking
- **Feature Engineer:** Domain-specific transformations
- **Model Manager:** Ensemble model orchestration
- **Interpretation Engine:** SHAP-based explainable predictions
- **Recommendation Generator:** Personalized health advice

## Frontend Application

### Framework and Components

- **Framework:** Streamlit 1.25+ for rapid ML application development
- **Styling:** Custom CSS for professional medical interface
- **Visualization:** Plotly for interactive charts and risk gauges
- **Responsiveness:** Mobile-optimized design patterns
- **State Management:** Session state for multi-page navigation

### User Interface Architecture



## Infrastructure and Deployment

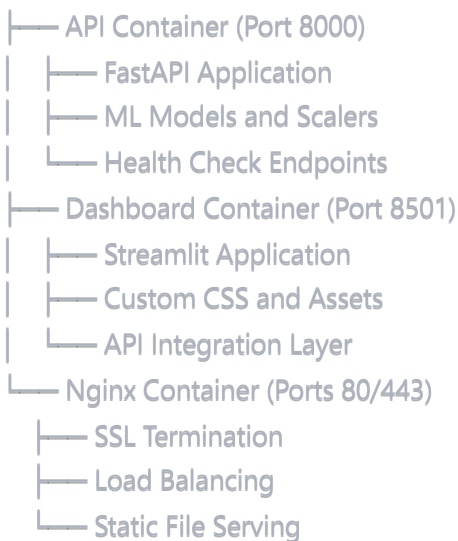
### Containerization Strategy

- **Multi-Service Architecture:** Separate containers for API, Dashboard, Proxy
- **Docker Images:** Optimized Python 3.9 slim base images
- **Build Strategy:** Multi-stage builds for production optimization
- **Volume Management:** Persistent storage for models and logs

### Container Architecture

yaml

#### Services:



### Orchestration

- **Docker Compose:** Multi-service orchestration for development and production
- **Service Dependencies:** Proper startup ordering and health checks

- **Network Isolation:** Custom bridge networks for service communication
- **Volume Persistence:** Data and model persistence across container restarts

## Production Deployment

- **Platform:** Ubuntu 22.04 LTS on Contabo VPS
- **SSL/TLS:** Let's Encrypt automatic certificate management
- **Domain Configuration:** heartdisease.duminduthushan.com
- **Monitoring:** Docker health checks and structured logging

## Data Flow Architecture

### Request Processing Pipeline

mermaid

graph TD

```
A[Patient Data Input] --> B[Input Validation Layer]
B --> C{Validation Passed?}
C -->|No| D[Error Response]
C -->|Yes| E[Feature Engineering Pipeline]
E --> F[Medical Domain Transformations]
F --> G[Feature Scaling and Normalization]
G --> H[ML Ensemble Models]
H --> I[Prediction Aggregation]
I --> J[SHAP Interpretation Engine]
J --> K[Risk Stratification Logic]
K --> L[Medical Recommendations Generator]
L --> M[Response Formatting]
M --> N[Client Application]
```

### Feature Engineering Pipeline

1. **Data Validation:** Medical parameter range checking
2. **Missing Value Handling:** Imputation strategies for clinical data
3. **Domain Feature Creation:**
  - Age stratification groups
  - Blood pressure categories
  - Cholesterol risk levels
4. **Interaction Features:**

- Age × Cholesterol interaction
- Blood pressure × Age interaction

#### 5. **Composite Metrics:**

- Heart rate reserve calculation
- Cardiovascular risk scoring

#### 6. **Feature Scaling:** RobustScaler for outlier handling

### **Model Inference Pipeline**

1. **Feature Preprocessing:** Apply trained scalers and transformations
2. **Ensemble Prediction:** Aggregate predictions from multiple models
3. **Probability Calibration:** Ensure reliable probability estimates
4. **Risk Classification:** Convert probabilities to risk categories
5. **Interpretation Generation:** SHAP-based feature importance
6. **Recommendation Engine:** Generate personalized health advice

### **Security Architecture**

#### **Data Protection**

- **Encryption in Transit:** TLS 1.3 for all client-server communications
- **Data Retention:** No persistent storage of patient health information
- **Input Sanitization:** Comprehensive parameter validation and sanitization
- **Rate Limiting:** Request throttling to prevent abuse (100 req/min)

#### **Access Control**

- **Network Security:** Container network isolation
- **Service Authentication:** Inter-service communication security
- **CORS Configuration:** Cross-origin resource sharing controls
- **SSL Configuration:** Strong cipher suites and security headers

### **Compliance Considerations**

- **HIPAA Readiness:** Architecture supports PHI handling requirements
- **Data Processing Transparency:** Clear data flow documentation
- **Audit Capability:** Comprehensive logging for compliance reporting

- **Privacy by Design:** Minimal data collection and processing

## Performance Architecture

### Scalability Design

- **Horizontal Scaling:** Stateless services for easy horizontal scaling
- **Load Distribution:** Nginx reverse proxy with upstream load balancing
- **Resource Optimization:** Efficient memory usage and CPU utilization
- **Caching Strategy:** Model and configuration caching for performance

### Performance Specifications

- **Response Time Targets:**
  - Single prediction: <2 seconds (95th percentile)
  - Batch processing: <1 second per patient
  - Dashboard rendering: <3 seconds initial load
  - API documentation: <1 second response
- **Throughput Capacity:**
  - Concurrent users: 100+ simultaneous assessments
  - Daily predictions: 10,000+ risk assessments
  - Batch processing: Up to 100 patients per request
- **Resource Requirements:**
  - Memory usage: <2GB per service container
  - CPU utilization: <80% under normal load
  - Storage requirements: <10GB total system storage

### Monitoring and Observability

- **Health Monitoring:** Comprehensive health check endpoints
- **Performance Metrics:** Response time and throughput monitoring
- **Error Tracking:** Structured error logging and alerting
- **Resource Monitoring:** Container resource usage tracking

# Deployment Architecture

## Environment Configuration

```
yaml

Production Environment:
  — Operating System: Ubuntu 22.04 LTS
  — Container Runtime: Docker 24.0+
  — Orchestration: Docker Compose
  — Web Server: Nginx 1.18+
  — SSL Management: Certbot/Let's Encrypt
  — Domain: heartdisease.duminduthushan.com
```

## Service Discovery and Communication

- **Internal Communication:** Container names for service resolution
- **External Access:** Nginx proxy with SSL termination
- **Health Checks:** Automated service health monitoring
- **Graceful Degradation:** Fallback mechanisms for service failures

## Backup and Recovery

- **Model Artifacts:** Regular backup of trained models and scalers
- **Configuration Backup:** Version-controlled configuration management
- **Container Image Storage:** Versioned container images for rollback
- **Recovery Procedures:** Documented disaster recovery processes

## Development and Operations

### CI/CD Pipeline (Ready for Implementation)

```
mermaid

graph LR
  A[Code Commit] --> B[Automated Testing]
  B --> C[Container Building]
  C --> D[Security Scanning]
  D --> E[Staging Deployment]
  E --> F[Integration Testing]
  F --> G[Production Deployment]
  G --> H[Health Monitoring]
```



## Quality Assurance

- **Automated Testing:** Unit, integration, and end-to-end test suites
- **Code Quality:** Linting, formatting, and static analysis
- **Security Scanning:** Container vulnerability assessment
- **Performance Testing:** Load testing and benchmarking

## Maintenance and Updates

- **Rolling Updates:** Zero-downtime deployment strategy
- **Configuration Management:** Environment-specific configurations
- **Logging Strategy:** Centralized logging with log rotation
- **Monitoring Alerts:** Proactive issue detection and notification

This architecture provides a robust, scalable, and maintainable foundation for the Heart Disease Risk Assessment System, demonstrating enterprise-grade software engineering practices suitable for production healthcare environments.